

[Cooperative Group Optimization] <http://www.wiomax.com/optimization>

Multiagent Optimization System for Solving the Traveling Salesman Problem (TSP)

Xiao-Feng Xie, *Member, IEEE*, and Jiming Liu, *Senior Member, IEEE*

Abstract—The multiagent optimization system (MAOS) is a nature-inspired method, which supports cooperative search by the self-organization of a group of compact agents situated in an environment with certain sharing public knowledge. Moreover, each agent in MAOS is an autonomous entity with personal declarative memory and behavioral components. In this paper, MAOS is refined for solving the traveling salesman problem (TSP), which is a classic hard computational problem. Based on a simplified MAOS version, in which each agent manipulates on extremely limited declarative knowledge, some simple and efficient components for solving TSP, including two improving heuristics based on a generalized edge assembly recombination, are implemented. Compared with metaheuristics in adaptive memory programming, MAOS is particularly suitable for supporting cooperative search. The experimental results on two TSP benchmark data sets show that MAOS is competitive as compared with some state-of-the-art algorithms, including the Lin–Kernighan–Helsgaun, IBGLK, PHGA, etc., although MAOS does not use any explicit local search during the runtime. The contributions of MAOS components are investigated. It indicates that certain clues can be positive for making suitable selections before time-consuming computation. More importantly, it shows that the cooperative search of agents can achieve an overall good performance with a macro rule in the switch mode, which deploys certain alternate search rules with the offline performance in negative correlations. Using simple alternate rules may prevent the high difficulty of seeking an omnipotent rule that is efficient for a large data set.

Index Terms—Cooperative systems, multiagent systems, optimization methods, traveling salesman problems (TSPs).

I. INTRODUCTION

THE TRAVELING salesman problem (TSP) [1], [2] is a classic combinatorial optimization problem. Although it can be easily formulated, it exhibits various interesting aspects of hard computational problems and has often served as a touchstone for novel approaches [3]–[6]. Moreover, TSP has various applications, such as very large scale integration (VLSI) design [7], rearrangement clustering [8], predicting protein functions [9], etc.

Due to the *NP-hardness* in a strong sense [10], various heuristic procedures [1], [2], [11], such as construction and local search (LS) strategies, have been applied for solving TSP instances with near-optimal solutions in reasonable amounts of running time. Representative LS strategies include 2-, 3-

[12], and 5-opt [13], the Lin–Kernighan heuristic (LK) [12], etc. Some simple metaheuristics, such as restart [14], [15] and iterated LS (ILS) [16]–[19], are also widely used.

Recombination search (XS) has been first introduced in a genetic algorithm (GA) [20]. For TSP, domain-specific XS strategies have been proposed from position- and order-based [21]–[23] to edge-based variants [24]–[29].

Various metaheuristics have been proposed, which may use common search strategies, such as construction, LS, XS, etc., as their components for solving TSP. Typical examples include ILS [16]–[19], GAs [20], [30], ant colony optimization (ACO) [31], [32], particle swarm optimization (PSO) [33], [34], etc.

Most state-of-the-art algorithms use certain LS heuristics, particularly LK variants, as the major search components during the runtime. For example, the LK–Helsgaun (LKH) [15] itself is an LK variant. In genetic LS (GLS) [27], [30], [35], [36], XS strategies are chained with LS strategies for producing high-quality solutions. Hence, among existing XS strategies, edge assembly crossover (EAX) [28], [37] attracts our interest since it had demonstrated its efficiency in certain selected TSP instances without resorting to the power of LS.

The multiagent optimization system (MAOS) [38] is a nature-inspired method, which addresses the self-organization [39] of agents working with limited declarative knowledge [40] and simple procedural knowledge under ecological rationality [41]. Specifically, agents explore in parallel [33], [42], based on socially biased individual learning (SBIL) [43], [44] and indirectly interact with other agents through sharing public information [45] organized in the environment (ENV) [46], [47]. Recently, MAOS has been applied to some hard problems [38], [48]. For TSP, the experiments on MAOS with an EAX variant have produced some primary encouraging results [48].

In this paper, MAOS is refined to implement simple and efficient knowledge components for solving TSP. In Section II, a simplified MAOS with limited memory specifications and simplified rules is described in a formalized form. In Section III, some existing knowledge components for TSP, particularly search behaviors and related auxiliary data structures, are described. In Section IV, MAOS is implemented with simple and efficient components for solving TSP. In Section V, the experimental results by MAOS cases on some TSP instances are presented and compared with some existing algorithms. In the last section, this paper is concluded.

II. MAOS

In this section, a multiagent optimization framework, in symbolic representation, is described. Afterward, a simplified MAOS version is defined based on the framework.

Manuscript received April 6, 2008; revised July 14, 2008. This paper was recommended by Associate Editor Y. S. Ong.

X.-F. Xie was with the Department of Computer Science, Hong Kong Baptist University, Kowloon Tong, Hong Kong. (e-mail: xiexf@ieee.org).

J. Liu is with the Department of Computer Science, Hong Kong Baptist University, Kowloon Tong, Hong Kong (e-mail: jiming@comp.hkbu.edu.hk).

Digital Object Identifier 10.1109/TSMCB.2008.2006910

A. Symbolic Representation

The general problem-solving capability arises from the interaction of declarative and procedural knowledge [49], [50].

Basic declarative and procedural knowledge components are represented in units called *chunks* and *rules*, respectively [49]. A chunk aggregates a small amount of information, and a rule specifies a certain action on certain information.

Moreover, macro components or modules may be organized by a combination of elemental knowledge components.

Each knowledge component is an object that can be defined by its *actual type* and certain *setting parameters*. Moreover, each action in a rule or a macro component has an interface with *input/output parameters*. Each component or parameter has a name. Each parameter has a *formal type*. Each *type* [51] is used for designating a specific organization structure.

For any two types, one type is defined as *compatible* with another type if it is an instance, i.e., a child type, of another type. Hence, the actual type of each valid parameter value must be compatible with the formal type of the parameter. Two types are not considered as compatible if they are different, unless there is an explicit declaration.

For a chunk, its actual type is a specific data structure. For a rule (R), its actual type is represented as $R_{I_KEY}^{I_NAME}$ [38], where the tags, i.e., I_KEY and I_NAME , designate the high-level interface with specific formal input/output parameters and the low-level realization, respectively. Any rules with a same I_KEY are compatible with a parameter in the formal type R_{I_KEY} . Here, R_{I_KEY} is also called an abstract rule.

For each module, the relations between its elemental components and their input/output parameters are hard-coded. Hence, a whole system can be symbolically described by all its setting parameters in scripts, where each actual parameter value may be an instance of a knowledge component or a value in a primitive data type. However, the number of setting parameters in such a system is not necessarily high since many components and parameters can be fixed in default values.

Such a system has the following features: 1) It allows for reusing and understanding of existing knowledge components; 2) the system may be locally improved by preferring those nondominated components with simplicity and efficiency, and 3) some knowledge components can be incrementally embedded into the system by using certain macro modes.

B. Multiagent Framework

The multiagent framework consists of a group of N_P autonomous agents. Each agent has a compact solving ability. The agents cooperatively search in an ENV [47], [48] for achieving a common intention of finding high-quality solution(s), based on the internal representation (IR) [48] of an optimization task, and related world knowledge [50].

The framework runs in iterative learning cycles. By running as a Markov chain process, the system behavior in the t th cycle only depends on the system status in the $(t - 1)$ th cycle. For convenience, during a run, $t = 0$ and $t > 0$ are called the *initialization* and the *runtime* stages, respectively.

Moreover, a chunk is called a *runtime chunk* if it is long-term-stored, retrieved and updated during the runtime. Each runtime

chunk aggregates certain particularities of potential solutions [52] of IR [48]. All rules and modules that are applied on the runtime chunks are also called behaviors.

1) *Problem Representation*: The IR [48] of an optimization task encapsulates certain basic knowledge of the task, which will be further processed for making solution(s) available [50]. Specifically, IR consists of a *landscape* for providing global structural information and certain *auxiliary knowledge components* associated with local structural information.

The landscape paradigm has been used for search in general [53]. Formally, a landscape is represented by a tuple, i.e., $\langle S_R, R_M \rangle$. S_R is called the *representation space*, containing all potential solutions to be searched, where each is called a *state* (\vec{x}). The *quality-measuring* rule (R_M) measures which one has a better quality between any two valid states in the S_R .

For a normal task, there is a cost function $f_C(\vec{x})$, which is to be minimized, for every state $\vec{x} \in S_R$. Then, the quality can be measured by a simple R_M^O rule [48], where $\forall \vec{x}_a, \vec{x}_b \in S_R$, if there is $f_C(\vec{x}_a) \leq f_C(\vec{x}_b)$, then \vec{x}_a has a better quality than \vec{x}_b , and R_M^O returns TRUE; else, it returns FALSE.

2) *Agent*: Each agent is a socially situated autonomous entity [42], [54]. Here, autonomy [54] indicates that an entity is able to control its own declarative and procedural knowledge components. Moreover, each agent is capable of generating new promising states, subjects to the limitations of available knowledge.

Each agent has several basic characteristics, as follows.

First, each agent possesses a personal long-term memory (LTM), called M_A . The memory stores some private runtime chunks. For an entity, each chunk in its LTM is persistent during its lifetime unless the memory is modified by itself. Most real-world animals possess LTM. The personal LTM supports *individual learning* [43], [44], which allows each entity to acquire knowledge during its lifetime by utilizing its past experience. In a group, individual learning allows agents to explore, in parallel, novel experiences in a large diversity.

Second, each agent refers to runtime chunks in public knowledge (M_S), which is organized in an ENV from the outside world of the agent [46]. Many species, including human beings [55], have evolved the capability for enhancing their group fitness [45] through *social learning* [56], [57], i.e., obtaining adaptive knowledge from their communities. Social learning also has its ecological significance for circumventing the inhibitory effects of individual learning [44].

Third, each agent can affect the outside world [46], which is realized by simply sharing a part of its knowledge to the ENV, i.e., at least one of the chunks in its M_A is nonprivate [45] and can be accessed by the outside world in the read-only mode.

Finally, the law of behavior is SBIL [43], [44] or cultural learning [58], i.e., a mix of reinforced practices of lifetime experiences and available public information [45], which has been adopted by many species for adaptive foraging [44]. As one of the *fast-and-frugal heuristics* in ecological rationality [41], it presents the following: 1) gains most of the advantages of both individual and social learnings [56] and 2) boosts collective properties by allowing adaptive experiences to be accumulated over many generations [56].

The solving capability is owned by a rudimentary central executive (CE) [48], [59] module. Specifically, CE performs a *generate-and-test* behavior [42], [60] for searching promising states, which contains a *generating* part, a *testing* part, and a *solution-extracting* part, under the support of a buffer memory, called M_G , which will be cleaned at the end of each cycle. Under the law of SBIL, the generating part creates new chunks and stores it into M_G by using the knowledge in both M_A and M_S . Afterward, the testing part updates M_A by using the chunks in M_G . The solution-extracting part just extracts all state(s) in M_G , which has no influence on the efficiency of the problem-solving process. Moreover, it is assumed that the generating part performs a rather deliberative behavior, while the testing part only produces a rather reactive behavior.

3) *ENV*: The ENV [46], [47] plays significant roles for supporting the solving capability of all agents.

First, ENV maintains all available world knowledge [50] and physical supports [47]. For example, it keeps a central clock that helps synchronize the behaviors, if necessary.

Second, the ENV serves as a task-dependent domain, namely: 1) It holds the IR of the optimization task, which encapsulates all problem-dependent knowledge, and 2) it recognizes candidate solution(s) for the task, which is simply realized by storing a quasi-solution state (\vec{x}^*), with the best quality among all promising states generated by agents for each single-objective optimization problem.

Finally, ENV also manages resources and services [47] for all agents according to their protocols. First, it provides the initial chunks in M_A of all agents by a *memory-initialization* module so that each agent itself can be designed in a simple way. Second, it organizes the corresponding M_S for every agent according to all available public knowledge [45], [61] by an interactive center (IC), through a protocol with all agents.

4) *Interaction*: Here, each agent may only directly interact with IC, in which the interaction protocol between them is based on an interface of two read-only information flows, i.e., 1) the nonprivate knowledge of the agent shared to IC, and 2) the corresponding M_S of the agent referred from IC. In principle, each agent and IC can be designed independently once the corresponding protocol is defined. It also implies that agents may be realized in a heterogeneous way, while the interaction protocol should be defined in a way that the realization of IC is not too complicated, particularly as there are multiple agents.

During the runtime, IC uses the nonprivate knowledge shared by all agents to update its private runtime chunks, and such knowledge will be organized as the M_S of each agent. Thus, agents indirectly interact with their peers through the medium role of IC. In fact, IC may act as a social memory [62], [63], where agents can share their past experiences for achieving emergent behavior by the self-organization of them [39].

C. Simplified System

Fig. 1 shows a simplified MAOS, which includes one of the agents and the ENV it roams. The agents are homogeneous in the sense that they have the same internal structure, which

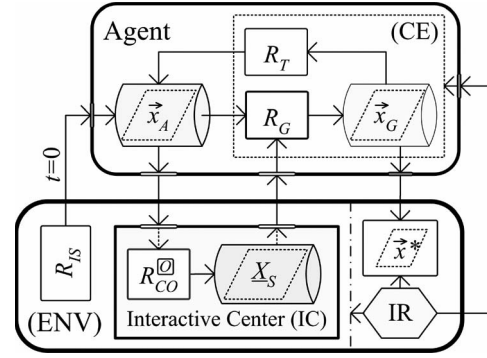


Fig. 1. One of the agents and the ENV it roams.

means that all agents are interacted with the ENV in the same way.

Given an IR, the chunks in three memories, including M_A , M_G , and M_S , and, afterward, the nonprivate chunk(s) in M_A , are specified in advance. Then, other components of MAOS may be implemented in a rather independent way.

1) *Basic Knowledge Components*: As an extremely limited version, both M_A and M_G contain only one chunk for holding the states \vec{x}_A and \vec{x}_G , respectively, and M_S contains only one chunk, i.e., the state set \underline{X}_S , which contains a set of states in S_R . Then, \vec{x}_A is exactly the only nonprivate knowledge to be accessed by IC.

The IC is designed in a centralized way [48]. It summarizes a state set, i.e., $\underline{X}_S = \{\vec{x}_{A(i)} | i \in [1, N_P]_{\mathbb{Z}}\}$, where each $\vec{x}_{A(i)}$ is referred from the M_A of the i th agent, by a basic *coordinating* rule (R_{CO}^O), and then, the same \underline{X}_S is shared by all the agents.

The memory-initialization module simply generates totally N_P states in the M_A of agents. Here, it is reduced to an *initial search* rule (R_{IS}), which generates each state independently.

The solution-extracting part is reduced as a trivial rule, which simply exports each \vec{x}_G directly.

The generating rule (R_G) is the major search rule during $t > 0$. It generates \vec{x}_G by using two input sources, i.e., \vec{x}_A and \underline{X}_S .

The testing rule (R_T) simply replaces \vec{x}_A by the generated \vec{x}_G once a specific criterion is satisfied. Due to its simplicity, the R_T may determine nontrivial properties of the $\vec{x}_A^{(t)}$ [48].

There are two simple R_T versions: 1) the *direct* $R_T(R_T^D)$, which does the replacement directly, and b) the *better* $R_T(R_T^G)$, which does the replacement as $R_M(\vec{x}_G^{(t)}, \vec{x}_A^{(t)}) \equiv \text{TRUE}$. Hence, $\vec{x}_A^{(t)}$ stores the most recent and the best states ever obtained by an agent as R_T^D and R_T^G are used, respectively.

The usage of *search rules* is a direct way of obtaining potential solutions. From a narrow perspective, a search rule is defined as a rule outputting a potential new state in S_R .

Both R_{IS} and R_G can be considered as search rules working during the runtime, respectively. However, R_T is not a search rule because it does not produce any new state.

The R_G rule may be decomposed into a filter and a search rule, where the filter may serve as a simple knowledge lens [64] for the purpose of organizing the desired declarative knowledge while suppressing the irrelevant.

A macro mode is called *generic* if it can embed with some *kernel* and auxiliary rules, and the type of such a macro rule is the intersection of the types of all kernel rules.

As a macro search rule, each kernel rule is a search rule. With macro search modes, we may only concern with those basic search rules. Then, various advanced macro search rules can be developed by applying different macro modes multiple times to embed existing search rules into them.

2) *Execution Process*: The simplified MAOS version is executed as follows.

At the initialization stage ($t = 0$), the states in M_A of all agents are to be supplied by the R_{IS} rule in the ENV, i.e.,

$$IR \xrightarrow{R_{IS}} \left\{ \vec{x}_{A(i)}^{(0)} | i \in [1, N_P]_{\mathbb{Z}} \right\} \quad (1)$$

where $\vec{x}_{A(i)}^{(0)}$ is the \vec{x}_A of the i th agent at $t = 0$.

During the runtime stage ($t > 0$), each cycle contains two sequential *clock steps* [48], i.e., C_PRE and C_RUN steps.

At the C_PRE step, the ENV itself is activated for executing some preparation operations by its components, if necessary. Specifically, the R_{CO}^O in IC is executed to organize the $\underline{X}_S^{(t)}$.

Then, at the C_RUN step, every agent is activated once for executing its rules. For the i th agent, it performs that

$$\vec{x}_{A(i)}^{(t)}, \underline{X}_S^{(t)} \xrightarrow{R_G} \vec{x}_{G(i)}^{(t)} \xrightarrow{R_T} \vec{x}_{A(i)}^{(t+1)}. \quad (2)$$

For each exported \vec{x}_G , if there is $R_M^O(\vec{x}_G, \vec{x}^*) \equiv \text{TRUE}$, then the \vec{x}_G is recognized as the quasi-solution state, i.e., \vec{x}^* .

The system is terminated if an optimal solution is found or if a specific overall cutoff criterion (R_{CCO}) is satisfied.

In summary, the simplified version has one overall setting parameter, i.e., N_P , and four rules, i.e., R_{IS} , R_G , R_T , and R_{CCO} . Each rule may also have certain setting parameters. Moreover, some rules may be decomposed into suitable macro forms in order to utilize some existing knowledge components.

D. Related Works and Advantage of MAOS

Here, we compare the current MAOS to some related works. Then, we have a discussion on the advantage of MAOS.

1) *Previous MAOS Versions*: The current framework is developed based on both the compact [38] and extended [48] versions. Since the compact version, a multiagent framework with a rudimental symbolic representation is formed. In the extended version, there are four changes: 1) IR is introduced for general search, including TSP; 2) clock steps are used for synchronizing behaviors in each clock; 3) a CE is taken into account for situating the generate-and-test behavior during the runtime; and 4) a fixed memory specification is applied for avoiding the complex management on declarative knowledge elements, as well as supporting an efficient search strategy.

In the current framework, there are four refinements: 1) The symbolic representation is associated with type system for easy reusing and understanding of the knowledge components; 2) a memory-initialization module is formally used, since the initialization may also have a nontrivial effect on the solving process; 3) an IC is formally introduced for managing the public information so that IC and agents can be flexibly designed according to an interaction protocol between them; and 4) a buffer memory (M_G) is defined for formally separating the three parts of the generate-and-test behavior.

In the simplified version, M_G only contains one state; thus, the testing behavior is simpler than both the previous versions, and the generating behavior itself simply becomes a search rule. Moreover, using generic macro search rules to embed with some basic search rule is particularly stressed.

2) *Population-Based Algorithms*: GA [20], ACO [31], [32], and PSO [33], [34] are three major examples.

GA simulates the evolution process. Each *individual* is a basic entity only containing a chromosome, which is not operated by the individual itself. Apparently, a *population* of these chromosomes forms the public information. Then, “the invisible hand” manipulates the population by using some evolutionary operators, such as selection, crossover, mutation, inheritance, etc. Instead, MAOS simulates the cultural learning process by a group of autonomous agents. Moreover, the simplified MAOS may directly use a rich pool of existing search operators proposed by GA researchers, if necessary.

ACO simulates the stigmergy mechanism in ant colonies. Each *ant* can be seen as an agent with simple behaviors; however, it differs from the agent in MAOS in that it does not possess a personal declarative memory. The *pheromone trails* laid down by the ants serve as the public information and can be implemented in the IC of MAOS with a suitable protocol. Without memory, an ant may only perform reflex behaviors. It may be interesting to know if any efficient strategies may be introduced once each ant possesses its personal memory.

PSO is a swarm-intelligence-based method. It has been implemented into the compact MAOS version [38], in which each *particle* is an agent. Each particle uses three chunks in its personal memory, and these chunks are associated with the search behavior based on a social-psychological principle. However, PSO is not supported in the simplified MAOS since it is still difficult to find competitive PSO variants for solving TSP.

3) *Advantage of MAOS for Cooperative Search*: As a simple multiagent system, MAOS may possess the potentials of parallelism, robustness, scalability, etc. [65]. Moreover, it is a bottom-up way to study social strategies [66]. However, here, we are only concerned with the advantage of MAOS in facilitating cooperative searches as compared with existing metaheuristics in adaptive memory programming (AMP) [52].

Cooperative search [67]–[69] consists of a search performed by multiple algorithms which share certain chunks during the runtime. With cooperation, statistical correlations among the performance of the individual algorithms will be introduced, which helps the cooperative portfolio, particularly as the offline performance of the individual algorithms are in negative correlations [67], to achieve better results than when the individual algorithms are independent of each other [67].

Many existing metaheuristics, including ILS, ACO, GA, etc., can be unified under the name of AMP [52] or a more sophisticated version called multiagent metaheuristic architecture (MAGMA) [68]. It should be mentioned that an agent in MAGMA could be actually reduced to a search behavior since it does not possess a personal memory. With a symbolic representation, AMP may be seen as rule-based heuristics. However, it is difficult to represent MAOS in AMP since AMP lacks a way of distinguishing the owners of runtime chunks.

In AMP/MAGMA, all the runtime chunks are stored as public information, and they have the primary importance to all search behaviors. It means that AMP/MAGMA only supports the cooperative search of its behaviors on public information. Thus, it should be very careful in adding/removing a behavior, since the change may have a significant impact on the whole system.

In MAOS, for each agent with autonomy, the runtime chunks from its personal memory and from the public information in IC have the primary and secondary importance, respectively. Thus, cooperative search may emerge in two different levels, i.e., the *agent* and the *group* levels.

For the agent level, the cooperation can be achieved if an agent deploys several alternate search rules, particularly in a macro realization of the generating part, during the runtime. From a perspective of multiple sequential learning cycles, the individual algorithms, i.e., the alternate rules, cooperate with others on the M_A of each agent in an interleaving way.

For the cooperation, the sharing information in the personal memory of an agent has the primary importance for the agent. For an agent in MAOS, the alternative search rules may be embedded independently for identifying, in advance, if their offline performances are in negative correlations. Moreover, the testing part may prevent certain high-risk results into its personal memory. Furthermore, the risk will be isolated in the agent only if an unexpected result occurs.

For the group level, the cooperation is applied on the group of agents. Each agent itself can be seen as an individual algorithm, and the agents cooperate with their peers only on public information in IC in a parallel way.

For the cooperation, the sharing information in IC has the secondary importance to each agent in the group. Thus, it does not have a great impact on the whole group, even as to add/remove heterogeneous agents. While the agents explore in parallel, the search process may be accelerated by their cooperation.

The group level is an inherent mechanism within MAOS. However, the agent level requires an agent using macro modules explicitly, as in the simplified MAOS version, which may be easily supported by the symbolic representation.

III. TSP DOMAIN KNOWLEDGE

The implementation of MAOS depends on the problem domain knowledge, which is described as follows. First, the IR of TSP is defined. Second, some auxiliary data structures, which are eligible for improving the search efficiency, are introduced. Third, some related search rules are described.

A. IR

For TSP [2], there is a graph with V nodes (or cities) and a cost matrix $D_O = (d_{jk})$, in which d_{jk} is the length of an edge that connects between cities j and k ($j, k \in [1, V]_{\mathbb{Z}}$). Here, we are only concerned with the symmetric TSP, which has $d_{jk} \equiv d_{kj}$ for the edges. The objective is to find a minimal-cost Hamiltonian tour, which passes through each node once and only once.

The representation space (S_R) contains all possible states, where each state \vec{x} represents a tour, which is a permutation $\{x_{(1)}, x_{(2)}, \dots, x_{(V)}\}$ of the integer values from 1 through V .

The cost function $f_C(\vec{x})$ to be minimized is represented as

$$f_C(\vec{x}) = d_{x_{(V)}x_{(1)}} + \sum_{j=1}^{V-1} d_{x_{(j)}x_{(j+1)}}. \quad (3)$$

Normally, the *distance* metric is defined as the number of different edges between any two tours [70].

The cost matrix D_O itself is the sole source associated with local structure information. For example, it can be used for building some candidate sets and for determining the relative cost changes for certain local moves, which are performed on valid tours, intermediate reference structures (REFs), etc.

B. Auxiliary Data Structures

Auxiliary data structures, e.g., reference structures (REFs) and candidate sets, are eligible for improving search efficiency.

1) *REF*: Some search strategies may use infeasible intermediate REFs, which are strongly associated with local structural information. For TSP, each REF may contain one or more disjointed graph elements in certain patterns, such as tour segment [25], [27], subcycle [11], [15], [28], stem-and-cycle (S&C) structure [71], etc.

Each REF can be modified into a valid tour by suitable reactive operations. Some graph elements can be preprocessed into simple elements. For example, each S&C structure can be transformed into a subcycle or segment. The REF in multiple tour segments (REF_{SM}) can be reconnected by randomly [29] or greedily [25], [27] inserting all missing edges. The REF in disjointed subcycles (REF_{CM}) can be turned into a valid tour by iteratively joining any two subcycles [15], [28].

2) *Candidate Set*: For TSP, a complete graph has totally $(V^2 - V)/2$ edges; however, most of the edges will not occur in good tours [2]. Thus, it is reasonable to focus on promising edges. One simple way is to use an edge set, which is called a *candidate set* (\underline{E}_C) [15], as a static sparse graph for restricting the search.

A simple example of \underline{E}_C is the nearest neighbor subgraph (NNS) [2], which keeps the nearest k_{nni} ($k_{nni} > 0$) neighbors for each node as candidates, based on the input matrix.

Usually, a minimum 1-tree shares many edges with an optimal tour. The length of the minimum 1-tree is often used as the lower bound [72], since it is a relaxation of TSP. In Lagrangian relaxation [15], [2], there is a penalized matrix $D_\pi = (d'_{jk})$, where $d'_{jk} = d_{jk} + \pi_j + \pi_k$, by associating each node j with a value π_j in a penalty array ($\vec{\pi}$). Intuitively, the transformed problem with D_π may be easily solved if it has a smaller length of the minimal 1-tree, which can be improved by applying a subgradient optimization [72] to $\vec{\pi}$. However, it might be rather time consuming to obtain a good $\vec{\pi}$ array [15], which requires one to solve minimum 1-trees many times, and each needs a time complexity $O(V^2)$ by Prim's algorithm.

Helsgaun have estimated the chances of a given edge being a member of a high-quality tour by using an α matrix [15]. The α matrix can be calculated [15] in a time complexity $O(V^2)$ and

a space complexity $O(V)$ based on a given input matrix, where each edge (j, k) has a corresponding α_{jk} value, which is equal to the minimum length as a 1-tree is required to contain this edge minus the length of a minimal 1-tree. Thus, it is intuitive that an edge is more promising to belong to an optimal tour if it has a smaller α value [15].

Two α matrices, i.e., D_α and $D_{\alpha\&\pi}$, are calculated by using D_O and D_π as the inputs, respectively. For example, the \underline{E}_C in the LKH [15] uses $D_{\alpha\&\pi}$ as its input matrix.

C. Types of Search Rules

Some types of search rules, which are frequency-used in existing algorithms, are defined based on the differences in their input parameters and internal features.

Scratch search rule (R_{SS}) generates a new state from scratch. Hence, it actually has an equivalent type with the R_{IS} rule.

As a child type of R_{SS} , *construction search* rule (R_{CS}) simply assigns the value of each part of a state once and only once.

Incumbent search rule uses one incumbent state as the input and searches its neighborhoods. It has two child types. The first is *perturbation search* (R_{PS}), also called *mutation* in GA [20] or *kick-move* in ILS [18]. The second is the LS rule (R_{LS}). The difference is that R_{LS} intends to improve the quality, which is capable of exploiting a landscape, while R_{PS} tries to move to a state with a distance away from the current state, which is useful for escaping from local optimum.

The XS rule (R_{XS}), which is similar to the *crossover* in GA [20], uses two parent states. An implicit advantage is that the distance between \vec{x}_O and each parent may be adaptively controlled by the distance of two parents.

Set-guided search rule (R_{SGS}) uses two parents; one is a state, and another is a state set. It has an equivalent type with the R_G rule. An example is an Inver-over operator [24].

All the aforementioned types can be easily embedded into the R_G rule with corresponding filters on their input information.

D. Generic Macro Search Rules

The *chained* macro mode has been used in some existing algorithms. Here, a macro rule contains a *chain* of two rules, i.e., a kernel and a R_{LS} rule, where a chain means that the input of a latter rule is exactly the output of a former rule.

Apparently, chained with the R_{LS} rule introduces a greedy search feature for the macro rule. A simple way of realizing a high-quality R_{SS} rule is to use an R_{CS} as the kernel rule [30]. In GLS [27], [30], [35], [36], the major search rule is a chained macro search rule by using an R_{XS} rule, which is normally not inherently greedy, as the kernel rule.

It may also increase the search capability if the kernel rule is an incumbent search rule, particularly as the R_{LS} rule does not allow uphill moves. For example, in ILS [18], the major search rule uses a perturbation search rule as the kernel rule.

E. Basic Search Rules

1) *Construction Search*: The construction search rule (R_{CS}) starts with a partial tour and adds the remaining nodes

one by one until the tour is complete. Here, two basic requirements are considered.

First, if there are multiple tours to be constructed, the diversity among these constructed tours should be large to provide enough information for search strategies.

The *randomized* R_{CS} rule (R_{CS}^R) is a simple R_{CS} version, which constructs a state \vec{x} in the S_R at random.

Second, each constructed tour should have a small distance to a good tour to contain positive information [15].

The *probabilistic* R_{CS} rule (R_{CS}^P) uses a candidate set ($\underline{E}_{C:CS}$) and D_O as its inputs. A probability matrix is initialized as $p_{ij} = d_{ij}^{-2} \forall i, j$. Then, each state is constructed by adding a city as the next city based on roulette selection from the remaining cities in ($\underline{E}_{C:CS}$) or else the remaining city with the maximum probability once all the cities in ($\underline{E}_{C:CS}$) were used.

2) *LS*: Well-studied LS strategies include r -opt variants, such as 2-, 3- [12], 5-opt [13], [15], etc., and multistage variants [10]. Some LS strategies may also use infeasible REFs, e.g., the two-stage moves with infeasible 2- [11], [13] and 3-opt [15], and the ejection chain algorithm [71] uses S&C for generating ejection moves.

As a simple LS, 3-opt systematically examines sequential exchanges (at most, three edges) in a positive gain criterion [12]. The full 3-opt version requires the examination of the $O(V^3)$ exchanges. In this paper, two 3-opt versions are considered. The R_{LS}^{3OC} rule simply adopts a static candidate set ($\underline{E}_{C:LS}$) for restricting the moves, and the R_{LS}^{3OS} rule incorporates one more speed-up technique, i.e., *don't look bits* [73] associated with each node.

3) *XS*: The XS rule (R_{XS}) generates a state \vec{x}_O by using two parent states, i.e., \vec{x}_{P1} and \vec{x}_{P2} .

Each tour is a permutation, as in some other problems, such as flow-shop scheduling [29], quadratic assignment problem [32], etc. Some *order-based* XS strategies try to preserve the order and position information, such as cycle crossover [23], partially mapped crossover [21], etc. Certain rules, e.g., maximal preservative crossover (MPX) [22], may also use edges as auxiliary information.

For TSP, each tour may serve as a restrict sparse graph [74]. In fact, edge information is more straightforwardly for guiding the search than order and position information [29].

Edge-based XS strategies try to utilize the positive edge information. For convenience, an edge is called *external* if it is not contained in the union edge set of both parents (\underline{E}_U). One basic idea is that any edge in (\underline{E}_U) can be freely introduced as the edge of \vec{x}_O . An extreme example is edge exchange crossover (EXX) [26], which does not add any external edges. Another basic idea is that good external edges should be also taken in account since (\underline{E}_U) does not necessarily contain all the edges belonging to, at least, an optimal tour, even as (\underline{E}_U) comes from many high-quality states [74]. Certainly, the quality of \vec{x}_O may be rather low if the probability of introducing unpromising external edges is high.

Many edge-based XS strategies work in two sequential steps:

- 1) selects a subset of the edges in (\underline{E}_U) into a REF, and
- 2) modifies the REF into \vec{x}_O . Only the second step may introduce certain external edges into \vec{x}_O .

Most existing strategies use the REF in multiple tour segments (REF_{SM}), although they may work under different classes, such as *intersection*, *union*, *partition*, etc.

The *intersection* class directly uses the intersection edge sets of both parents as the REF_{SM}. A famous example is distance-preserving crossover [27].

The *union* class uses the union edge sets of parents for constructing tour segments. In edge recombination [29], tour segments are constructed by considering node degrees.

The *partition* class is realized as follows. First, the nodes are partitioned into different groups. For each parent tour, each group is associated with a subgraph with corresponding edges connecting the nodes within the group. Second, the REF_{SM} is formed by a union set of all groups from different parents. A typical example is natural crossover [25].

A few edge-based XS strategies utilize the intermediated REF in disjointed subcycles (REF_{CM}). Most of them belong to the family of the EAX [28].

The probability of introducing low-quality edges by joining subcycles in REF_{CM} is much lower than that by connecting segments in REF_{SM}, as both kinds are realized in greedy mechanisms, since the former has more choices. Joining may be applied on all nodes in subcycles, while connecting can only happen on the two end nodes of each segment.

IV. IMPLEMENTATION

For convenience, certain template(s) may be specified in advance, and certain basic MAOS case(s) can be defined by implementing all rules; then, all related cases can be described by applying local modification(s) on the basic cases.

Two simple improving heuristics are taken into account, based on a generalized edge assembly recombination (GEAX) for enhancing the overall search performance.

A. GEAX

The GEAX is a generalization of EAX [28], [37], which is realized as follows.

Step 1) Design an auxiliary set, called *AB-Set*, where each is called an *AB-cycle* (\underline{E}_{AB}) [28], which is an even-length cycle with *A-Edges* (\underline{E}_A) and *B-Edges* (\underline{E}_B), by tracing different edges of \vec{x}_{P1} and \vec{x}_{P2} alternately. It is realized in two substeps.

Step 1.1) Build the *AB-library* by using \vec{x}_{P1} and \vec{x}_{P2} . As each AB-Cycle is identified, it is stored into the AB-Library, and the edges that compose it are no longer used [28]. This procedure is repeated until no AB-Cycle could be extracted or the number of AB-Cycles has achieved an upper value (N_{ABL}).

Step 1.2) Design the *AB-Set* by an *AB-selecting* rule (R_{SAB}), which selects some AB-Cycles from the AB-Library.

Step 2) Generate multiple candidate states by using both the \vec{x}_{P1} and the AB-Set, where each AB-Cycle in the AB-Set is used for guiding the kick-move on \vec{x}_{P1} in two substeps [28].

Step 2.1) Kick the \vec{x}_{P1} into a REF_{CM} structure, called \underline{E}_O , which contains disjointed subcycles covering all nodes, by using each AB-Cycle. First, \vec{x}_{P1} is copied as the initial \underline{E}_O . Then, all A and B-Edges in the AB-Cycle are removed from and added to the \underline{E}_O . Thus, all the edges introduced in the \underline{E}_O , i.e., the B-Edges, are selected from the \underline{E}_{P2} , based on the definition of the AB-Cycle.

Step 2.2) Modify the REF_{CM} structure into a valid tour. It is achieved by joining two subcycles iteratively, which is accelerated by a candidate set ($\underline{E}_{C:XS}$). While there are multiple subcycles, the subcycle with the minimal number of edges is selected as the base one, and then, a greedy two-exchange with maximum local gain is used for merging it with any other one, where an edge is systematically selected in the $\underline{E}_{C:XS}$.

Step 3) Choose one state among the candidate states as the \vec{x}_O by a *state-competing* rule, which is simply realized by associating each candidate state \vec{x}_c with the competing function (f_{comp}) to be minimized [48], i.e.,

$$f_{comp}(\vec{x}_c) = \frac{f_C(\vec{x}_c) - f_C(\vec{x}_{P1})}{DIS(\vec{x}_{P1}, \vec{x}_c)^{C_D}} \quad (4)$$

where $DIS(\vec{x}_{P1}, \vec{x}_c)$ reflects the distance between \vec{x}_{P1} and \vec{x}_c and C_D is the *diffusion coefficient*, which may determine the velocity of information diffusion [48]. The $DIS(\vec{x}_{P1}, \vec{x}_c)$ function is approximately estimated as $|\underline{E}_{AB}|/2$, i.e., a half of the number of edges in the corresponding AB-Cycle [37].

The f_{comp} integrates both the high-quality criterion [75] and the distance requirement [18], [27], [70]. If $C_D = 0$, it represents the high-quality criterion only, as in EAX-1AB [28] and some of its variants [75], [76]. As $C_D < 0$, it also prefers high-quality states that are far away from the incumbent state, as in the fitness-distance-based diversification [18]. As $C_D > 0$, it prefers high-quality states closer to the incumbent state. If $C_D = 1$, it represents the criterion employed in EAX-Dis [37]. In [48], $C_D = 1.5$ is used for slowing down information diffusion.

The inherent greedy feature is achieved by two aspects. First, \vec{x}_{P1} is kicked into REF_{CM} by introducing the edges in \vec{x}_{P2} (Step 2.1), while the REF_{CM} can be turned into a valid tour with a few promising external edges (Step 2.2). Second, there is an iterative child generation (ICG) [18], [75] (Step 2), i.e., the competition among multiple independent candidate trials. Such a simple competing algorithm portfolio may reduce the heavy-tailed distributions of individual trials [14], [67]. In fact, ICG has shown its efficiency in ILS [18], and it may perform better than 2-opt as it works with EAX [75].

In summary, GEAX has four variable parts, i.e., C_D , N_{ABL} , R_{SAB} , and $\underline{E}_{C:XS}$, where N_{ABL} is fixed as 100 in this paper.

The random GEAX (GEAX.R) uses the *random* R_{SAB} rule (R_{SAB}^R), where totally N_{MT} AB-Cycles are randomly

chosen from the AB-Library and any repeated AB-Cycles are discarded.

Some existing EAX variants can be represented by GEAX cases. For example, EAX-Dis [37] is the GEAX.R case with $C_D = 1$ if the potential difference in $\underline{E}_{C:XS}$ is not considered.

B. Simple Improving Heuristics

Under limited time and resources, animals make inferences about specific domains in the real world with simple heuristics by utilizing certain imprecise knowledge [41].

1) *Sorting-Based AB-Cycle Selection*: The first heuristic is applied into the process of ICG [18], [75] so as to utilize imprecise positive clues before generating candidate states.

For GEAX, some AB-Cycles are obtained in Step 1.1), where each AB-Cycle leads to a candidate state. However, the random R_{SAB} rule does not utilize any information available at the stage.

Each AB-Cycle (\underline{E}_{AB}) is associated with the function f_{clue}

$$f_{clue}(\underline{E}_{AB}) = \sum \text{len}(\underline{E}_B) - \sum \text{len}(\underline{E}_A) \quad (5)$$

where $\sum \text{len}(\underline{E})$ represents the sum of length of all edges in \underline{E} . Hence, the f_{clue} value is the delta length of total edges between \bar{s}_{P1} and the intermediate REF_{CM} structure to be generated by the \underline{E}_{AB} at Step 2.1). Intuitively, an AB-Cycle with a smaller f_{clue} value may carry an imprecise positive clue, although some new edges will be included later to modify the REF_{CM} into a valid tour at Step 2.2).

The *sorting* R_{SAB} rule (R_{SAB}^S) is proposed at Step 1.2), where totally N_{MS} AB-Cycles in the AB-Library with the minimal f_{clue} values are selected as the members of AB-Set. Certainly, if N_{MS} is larger than the size of the AB-Library, then all the elements in the AB-Library become the members of the AB-Set.

The *sorting* GEAX (GEAX.S) simply differs from GEAX.R in that the former uses R_{SAB}^S instead of R_{SAB}^R at Step 1.2) of GEAX.

2) *Switch Macro Mode*: The second heuristic is the generic macro search rule in *switch* mode. The switch macro mode takes multiple kernel rules in the same I_KEY as the alternate rules, which may have different I_NAME tags, and only runs one of them, which is selected by a *deploying rule* (R_{DEP}) [42], i.e., a simple task-switching schedule [77], at each activation.

If R_G or at least one of its components (e.g., R_{XS}) is realized in the switch mode, then agents may run alternative R_G rules at each cycle. For each agent, all the alternate R_G rules cooperate, at least, by its personal knowledge during multiple cycles. Thus, the switch macro mode supports the cooperative search of multiple alternative search rules in the agent level.

C. Standard Template

The standard MAOS template, which employs N_P compact agents, has the following specific components.

The R_{IS} rule is simply a macro scratch search rule in chained mode. It is organized as a tuple $\langle R_{CS}, R_{LS} \rangle$, i.e., a chain of a construction search rule (R_{CS}) and an LS rule (R_{LS}).

The R_G rule is decomposed as a tuple $\langle R_{SP}, R_{XS} \rangle$, which contains a *state-picking* rule (R_{SP}) and an XS rule (R_{XS}). For the R_{XS} rule, a basic requirement is that it is a greedy strategy. The R_{XS} rule uses two input sources, i.e., \bar{x}_{P1} and \bar{x}_{P2} , where $\bar{x}_{P1} = \bar{x}_A$ and \bar{x}_{P2} is filtered from \underline{X}_S by the R_{SP} rule. For the R_{XS} rule, \bar{x}_{P1} and \bar{x}_{P2} can be seen as the incumbent state and the guiding information, respectively.

Moreover, the macro search rule in *switch* mode is applied on the R_{XS} rule. The macro R_{XS} rule contains an R_{DEP} rule and several alternate R_{XS} rules. If there is only one alternate R_{XS} rule, then the macro R_{XS} is reduced to the alternate R_{XS} .

Organized by simple components with specific roles, the MAOS template tends to tailor to the “big valley” [70] in the TSP landscape, which suggests that better local minima tend to have a smaller distance to the closest optimum by sharing common partial structures [17], [19], i.e., edges or segments.

At the initialization stage, the R_{CS} is to construct states with both quality and diversity, and the R_{LS} is to improve the quality of each state for reaching the “big valley.” Then, during the runtime, as the basic intuition behind ILS [16]–[18], agents explore in parallel, where each high-quality state $\bar{x}_{A(i)}$ may serve as the incumbent state for searching global optimum in its neighborhood with a short distance.

Moreover, the high-quality tours in \underline{X}_S tend to concentrate on a very small subspace around the optimal tour(s) [32], which may lead to two positive effects: 1) The union of edges in high-quality tours may be seen as a restrict graph. The graph may cover most or even all edges in an optimal tour while it is quite sparse, due to the large number of shared edges among high-quality tours [74]. The sparse restrict graph may serve as a dynamic guiding information for facilitating the search process. It has been shown in ILS [16] that cost-restricted kicks can be more effective than blind kicks, particularly for large instances, and b) the edge set may represent pseudobackbone frequencies [19], although it is *NP hard* to find the exact backbone [78]. Compared with the pheromone matrix [31], [32], the edge set provides not only the frequency of edges but also the frequency of many tour segments, for belonging to high-quality tours. Through picking the partial information by the R_{SP} rule, it is able to utilize the pseudobackbone frequencies implicitly.

D. Standard Case

The standard case (#STD) is defined based on the standard template, where the N_P will be specified at each time it is used.

For R_{IS} , it is implemented as $\langle R_{CS}^P, R_{LS}^{3OS} \rangle$. For R_T , R_T^G is used. For the macro R_{XS} rule, the GEAX.R is employed as the alternate R_{XS} rule, which has $C_D = 1.0$ and R_{SAB}^R with $N_{MT} = 20$. Basically, GEAX.R is equivalent to EAX-Dis [37] if the difference in $\underline{E}_{C:XS}$ is not considered.

For R_{DEP} , the *randomized* R_{DEP} rule (R_{DEP}^R) is realized by selecting one of the alternate rules at random [42].

An identical candidate set ($\underline{E}_{C:STD}$) is employed for restricting the search of all related behavioral rules, i.e., R_{CS} , R_{LS} , and R_{XS} . Specifically, $\underline{E}_{C:STD}$ is realized in the NNS mode [2], which has $k_{nni} = 20$ and has D_O as its default input matrix.

For R_{SP} , the *randomized* R_{SP} rule (R_{SP}^R) is considered, which picks one of the states from the given \underline{X}_S at random.

For R_{CCO} , the R_{CCO}^{MC} rule is used, which is terminated if the maximum number of cycles (T_{MAX}) is achieved or if no further improvement on the solution quality occurs for T_{CON} cycles. Here, T_{MAX} and T_{CON} are fixed as 500 and 100, respectively.

For each tour, the most straightforward *Array* representation [79] is used by both the R_{CS} and R_{LS} rules, while the *doubly linked list* representation [79], which has a similar time complexity with the *Array* representation, is used by GEAX since it is convenient for representing the REF_{CM} structure.

V. EXPERIMENTAL RESULTS AND DISCUSSION

The MAOS is coded in JAVA, and was run by JRockit JVM 1.5 on a 1.5-GHz Itanium 2 processor with 1-G memory.

Two benchmark sets are used; the first is all 21 instances with $1000 \leq V < 3000$ nodes in the TSP library (TSPLIB) [80], and the second is all 22 instances with $3000 \leq V < 6000$ in the VLSI data set [7]. Using a set of instances with nodes in a range may bring a lesser bias on the average performance than by using some selected instances. For each instance, ten independent runs (N_R) were performed.

There are three indices for the performance of an algorithm. The first is the success rate (p_s), which is estimated by N_S/N_R , where N_S is the number of times optimal states were found. The second is the relative percentage deviation (RPD) above the best known solution (f^*). The third is the running time (t_r) in seconds, which may be influenced by different machine configurations. Both p_s and RPD are related to the solution quality, although they are based on different viewpoints.

A. Contributions of Components and Parameters

The contributions of the MAOS components are evaluated on the TSPLIB instances by using RPD–Time diagrams for indicating the Pareto efficiency of both performance indices. For each case, $N_P = 10, 30, 50, 100$, and 300 are tested.

Fig. 2 shows the results by #STD using different R_T and R_{IS} components. For R_T , #STD|RT.D simply uses R_T^D instead of R_T^G , which removes the quality control. It produces similar performance with #STD, which means that the employed R_{XS} has an inherent greedy feature in a strong sense.

For initialization, the #STD|CS.RND uses R_{CS}^R instead of R_{CS}^P , and #STD|LS.3OC uses R_{LS}^{3OC} instead of R_{LS}^{3OS} . In fact, the tuple $\langle R_{CS}^R, R_{LS}^{3OC} \rangle$ has been used in [48] for initialization, where R_{LS}^{3OC} has shown its advantage as compared with the 2-opt [48]. Fig. 2 shows that #STD achieves highly dominating performance than both #STD|LS.3OC and #STD|CS.RND.

Table I lists the average information at $t = 0$ for all three cases under $N_P = 10, 100$, and 1000, where RPD_0 and t_{r0} are RPD and t_r at $t = 0$, respectively, and $r_{T0} = t_{r0}/t_r$. It can be seen that #STD achieves much lower r_{T0} than both of the others. It reduces both RPD_0 and t_{r0} significantly by using R_{CS}^R instead of R_{CS}^P , which indicates the high advantage of using suitable R_{CS} heuristics, even with simple input information. With R_{LS}^{3OS} , which simply incorporates *don't look bits* [73] into R_{LS}^{3OC} , the initialization requires much less t_{r0} , although it also produces worse RPD_0 . Certainly, it still leaves a large space for improving the R_{IS} since the initialization costs more than 32%




Fig. 2. Results by #STD using different R_T and R_{IS} rules.

TABLE I
AVERAGE RESULTS AT THE INITIALIZATION STAGE ($t = 0$)

Case	$N_P=10$			$N_P=100$			$N_P=1000$		
	RPD ₀	t_{r0} (s)	r_{T0}	RPD ₀	t_{r0} (s)	r_{T0}	RPD ₀	t_{r0} (s)	r_{T0}
#STD	5.61	0.55	0.41	4.44	3.84	0.32	3.85	35.70	0.34
#STD CS.RND	9.04	0.97	0.52	6.73	7.86	0.49	5.03	74.89	0.49
#STD LS.3OC	3.90	1.96	0.69	3.06	16.50	0.69	2.53	160.36	0.71

of the total running time. For investigating the performance of R_{XS} , it is expected that the R_{IS} to be introduced may reduce r_{T0} , while it is not too sophisticated.

Here, the t_{r0} is mainly consumed by the R_{LS} . The *Array* representation may be replaced by the *two-level tree* [79], which may drop the time complexity of per *Flip* from $O(V)$ to $O(\sqrt{V})$, while having no impact on its outputs. It has also been declared that the “segment list” version [30] may make it twofold faster than the conventional *two-level tree*. Moreover, more sophisticated R_{LS} may be considered, concerning the overall performance. For example, 5-opt [13] may achieve better performance than 3-opt, which has been the basic move in LKH [15] and PHGA [30]. The chained LK (CLK) [16], which has been used in some algorithms [36], may also be considered.

Figs. 3 and 4 show the results by using different N_{MT} values for R_{SAB}^R and different C_D values for GEAX.R, respectively.

Fig. 3 shows that the performance of #STD is quite stable for N_{MT} varied from 10 to 50 and only becomes worse as N_{MT} is smaller, such as $N_{MT} = 5, 3$, and 1. Hence, the choice of N_{MT} is quite flexible. Moreover, it also suggests the advantage of ICG [18], [75].

Fig. 4 shows that varying C_D from -1.0 to 2.0 does not introduce a great difference, although the case with $C_D = -1.0$ seems slightly worse, which may be due to two possible reasons. The first is that the multiagent framework may facilitate the preservation of the diversity of positive clues in public information [58], as explored by the agents in parallel, where each of them possesses a private memory and only picks a piece of public information for biasing its search. The second is that the GEAX worked in the biased style may keep a local diffusion effect, since most AB-Cycles are in small sizes and some edges may be brought back by the greedy joining mechanism. Aside




Fig. 3. Results by #STD using different N_{MT} values for the R_{SAB} rule.




Fig. 4. Results by #STD using different C_D values for the GEAX.R.

from this, a larger C_D value may be considered for achieving better performance if a longer running time is available.

Fig. 5 shows the results by using candidate sets ($E_{C:STD}$) in the NNS mode, which are derived under four input matrices, i.e., $D_O(\#STD)$, $D_\alpha(\#STD_\alpha)$, $D_\pi(\#STD_\pi)$, and $D_{\alpha\&\pi}(\#STD_{\alpha\&\pi})$, respectively. The penalty array ($\vec{\pi}$) for obtaining D_π of each TSBLIB instance comes from the LKH 1.3 [15]. It shows that using D_α , D_π , and $D_{\alpha\&\pi}$ as inputs of candidate sets may bring on a better performance than using D_O . However, it should be careful to use $\vec{\pi}$ by counting the high time complexity for finding a suitable one [15].

B. Performance Improvements

For further investigating the performance of MAOS cases with improving heuristics, $\#STD_\alpha$ is taken as a basic case.

The case $\#STD_\alpha|S$ simply uses GEAX.S, which has $N_{MS} = 10$ for the R_{SAB}^S rule, as the R_{XS} rule to be used.

The case $\#STD_\alpha|S\&R$ deploys two alternate R_{XS} rules, i.e., the GEAX.R in $\#STD_\alpha$ and the GEAX.S in $\#STD_\alpha|S$, by the R_{DEP}^R rule. It is equivalent to a macro R_G rule with two alternate R_G rules, where each uses one of the R_{XS} rules.




Fig. 5. Results by #STD using candidate sets in different input matrices.




Fig. 6. Results by #STD and $\#STD_\alpha$ in different R_{XS} 's, on TSPLIB instances.


Fig. 6 and 7 show the results of four MAOS cases, i.e., $\#STD$, $\#STD_\alpha$, $\#STD_\alpha|S$, and $\#STD_\alpha|S\&R$, which are applied on the TSPLIB and VLSI instances, respectively.

First, it shows that the $\#STD_\alpha|S$ performs quite well on both RPD and t_r for the TSPLIB instances. Hence, the sorting criterion by R_{SAB}^S may indeed catch certain positive clues for achieving fewer trials in ICG.

Second, it also shows that $\#STD_\alpha|S$ produces worse results for the VLSI data set as N_P is large. Hence, the MAOS cases with GEAX.R and GEAX.S are in negative correlations [67], where each is particularly good at one of the data sets.

Thirdly, it demonstrates the efficiency of $\#STD_\alpha|S\&R$ on both data sets. The $\#STD_\alpha|S\&R$ case achieves a similar performance with the better one of $\#STD_\alpha$ and $\#STD_\alpha|S$. As N_P is larger than 100, the cooperative search by GEAX.R and GEAX.S even produces better performance than the search by GEAX.R or GEAX.S alone for the VLSI data set.

Such a cooperative search may imply a different viewpoint for designing algorithms, since the offline performance of each

Fig. 7. Results by #STD and #STD $_{\alpha}$ in different R_{XS} 's, on VLSI data set.TABLE II
RUNNING TIMES OF THE GREEDY HEURISTIC [1] ON DIFFERENT CPUs

CPU\Time\V	10^3	$10^{3.5}$	10^4	$10^{4.5}$	10^5	$10^{5.5}$	r_{CPU}
Itanium 2 (1.5GHz)	5.39	5.70	6.58	8.45	9.89	15.47	1.00
Pentium III (933MHz) [30]	11	11	20	35	43	48	2.82
Alpha (500MHz) [1]	12	14	16	23	36	55	2.74

alternate rule can be accumulated during the solving process. Instead of seeking for an omnipotent rule for a full set of tasks, which might be extremely difficult, if not impossible, the macro rule for deploying simple alternate rules in negative correlations may be considered. Moreover, only a part of the new alternate rules need to be implemented for covering those new tasks that are difficult for existing alternate rules, as the domain structure features of tasks may change over time.

C. Comparisons With Other Algorithms

It is difficult to compare the running times of various algorithms on different executing conditions, including both machine configurations and programming platforms.

A simple way to roughly compare the difference on various machine configurations is to execute a same benchmark code [1]. Table II summarizes the running times of the C code of a greedy heuristic [1]. The r_{CPU} values mean that the 1.5-GHz Itanium 2 is about 2.82 and 2.72 times faster than the 933-MHz Pentium III [30] and the 500-MHz Alpha [1], respectively.

Moreover, both the JAVA and C benchmark codes of SciMark 2.0 were executed on the 1.5-GHz Itanium 2 by using JRockit 1.5 and GCC 3.3.3 (with $-O2$ option) and obtained the composite scores of 170.63 and 185.08 (higher is better), respectively. It indicates that the performance gap between JAVA and C is really close. For simplicity, the small gap is not considered in the comparisons.

For the two CPUs with r_{CPU} values in Table II, we compare the normalized running times. For other CPUs compared in this paper, we simply compare the real running times; however, the readers may also achieve a rough normalization since such CPUs are faster than the 933-MHz Pentium III in Table II.

TABLE III
RESULTS BY LKH, IBGLK, AND PHGA ON TSPLIB INSTANCES

Name	f^*	LKH			IBGLK		PHGA		
		p_s	RPD	t_r (s)	RPD	t_r (s)	p_s	RPD	t_r (s)
dsj1000	18659688	0.90	0.0001	88.2	0.010	2066	0.90	0.0001	2860.7
pr1002	259045	1.00	0.0000	25.9	0.000	345	1.00	0.0000	507.3
sil032	92650	0.40	0.0072	119.6	0.000	772	1.00	0.0000	1037.5
u1060	224094	0.80	0.0142	35.4	0.020	873	1.00	0.0000	321.2
vm1084	239297	0.80	0.0092	48.3	0.020	403	1.00	0.0000	1089.4
pcb1173	56892	0.90	0.0009	42.1	0.000	213	1.00	0.0000	419.2
d1291	50801	0.40	0.0669	134.1	0.079	1086	1.00	0.0000	813.1
rl1304	252948	0.70	0.0324	49.3	0.000	504	1.00	0.0000	651.5
rl1323	270199	0.10	0.0131	48.1	0.000	582	1.00	0.0000	1193.1
nrw1379	56638	0.10	0.0065	68.9	0.060	440	1.00	0.0000	870.0
fl1400	20127	0.00	0.1913	354.6	0.000	19317	1.00	0.0000	7544.3
u1432	152970	1.00	0.0000	114.4	0.000	617	1.00	0.0000	739.3
fl1577	22249	0.00	0.0539	659.8	0.000	10430	0.60	0.0090	8435.6
d1655	62128	1.00	0.0000	105.3	0.000	1869	1.00	0.0000	1258.9
vm1748	336556	0.60	0.0153	127.2	0.010	829	1.00	0.0000	1587.7
u1817	57201	0.00	0.0960	173.5	0.159	680	1.00	0.0000	1638.6
rl1889	316536	0.60	0.0011	119.2	0.217	1110	1.00	0.0000	1483.5
d2103	80450	0.00	0.0292	492.6	0.000	8669	0.90	0.0006	6841.3
u2152	64253	0.20	0.0408	304.9	0.189	808	1.00	0.0000	1411.6
u2319	234256	1.00	0.0000	957.5	0.040	2091	0.00	0.0350	3200.5
pr2392	378032	1.00	0.0000	275.0	0.138	715	1.00	0.0000	947.8
Average	-	0.55	0.0275	206.9	0.117	2679	0.92	0.0021	2135.8

The performance of MAOS on the TSPLIB instances are compared with three algorithms, i.e., LKH [15], IBGLK [19], and PHGA [30]. All the algorithms are coded in C. The times of IBGLK and PHGA were normalized to 500-MHz Alpha [19] and 933-MHz Pentium III [30], respectively, and the time of LKH was measured on a 933-MHz Pentium III [30].

The IBGLK [19] combines iterated LK algorithms with the pseudobackbone frequencies by considering two techniques: 1) construct pseudobackbone by using samples of local minima and 2) alternate between BGLK and regular LK once each of them fails. It has outperformed the three LK versions.

The LKH [15] has been one of the most powerful LK variants, which considers a powerful set of two-stage moves and utilizes $D_{\alpha \& \pi}$ as the input matrix of a candidate set. The results tested in [15] did not count the time for obtaining suitable $\bar{\pi}$ arrays and missed one of the instances, i.e., d2103. Hence, the results tested in [30], which were performed on the UNIX version of LKH 1.3 [15], are considered.

The PHGA [30] is a multipopulation hybrid GA, which uses a variant of the MPX [22].

Some speed-up techniques, such as candidate set, *don't look bits* [73], and the *two-level tree* [79] tour representation, are used. Moreover, PHGA represents tours in the "segment list" version [30] of the *two-level tree* and uses Helsgaun's 5-opt [15].

Table III summarizes the performance indices of LKH [30], IBGLK [19], and PHGA [30], for the TSPLIB instances.

Table IV lists the results of #STD $_{\alpha}$ |S&R with $N_P = 300$, 500, and 1000 for the TSPLIB instances, where the T_{RUN} means the actual running cycles. With $N_P = 1000$, MAOS can achieve $p_s > 0$ for all the instances and $p_s = 100\%$ for 19 instances.

When comparing Table IV with Table III, the MAOS with $N_P = 300$ achieves performance 1.44 times higher for the p_s , 6.71 times lower for the RPD, and 2.10 times faster for the normalized running time than that of LKH; the MAOS with

TABLE IV
RESULTS BY #STD_α|S&R WITH DIFFERENT N_P 's
ON TSPLIB INSTANCES

Name	$N_P=300$			$N_P=500$			$N_P=1000$		
	p_s	RPD	$t_r(s)$	p_s	RPD	$t_r(s)$	p_s	RPD	$t_r(s)$
dsj1000	1.00	0.0000	21.3	1.00	0.0000	35.6	1.00	0.0000	68.9
pr1002	1.00	0.0000	19.5	1.00	0.0000	32.3	1.00	0.0000	64.2
si1032	0.50	0.0005	9.4	0.70	0.0003	14.9	1.00	0.0000	28.8
u1060	1.00	0.0000	21.1	1.00	0.0000	34.9	1.00	0.0000	67.0
vm1084	1.00	0.0000	17.0	1.00	0.0000	27.2	1.00	0.0000	51.6
pcb1173	1.00	0.0000	21.4	1.00	0.0000	34.9	1.00	0.0000	69.1
d1291	1.00	0.0000	15.6	1.00	0.0000	25.4	1.00	0.0000	50.0
rl1304	1.00	0.0000	14.6	1.00	0.0000	23.9	1.00	0.0000	45.2
rl1323	1.00	0.0000	16.4	1.00	0.0000	26.0	1.00	0.0000	50.1
nrv1379	1.00	0.0000	44.1	0.90	0.0009	71.6	1.00	0.0000	144.0
fl1400	0.90	0.0010	26.2	1.00	0.0000	42.7	1.00	0.0000	82.1
u1432	0.80	0.0024	36.5	0.90	0.0012	60.4	1.00	0.0000	116.5
fl1577	0.20	0.0135	27.9	0.60	0.0063	46.0	0.50	0.0067	92.8
d1655	0.70	0.0026	34.8	1.00	0.0000	56.6	1.00	0.0000	112.7
vm1748	1.00	0.0000	40.8	1.00	0.0000	66.3	1.00	0.0000	131.8
u1817	0.40	0.0285	51.1	0.90	0.0037	83.4	1.00	0.0000	156.4
rl1889	1.00	0.0000	35.5	1.00	0.0000	58.1	1.00	0.0000	115.7
d2103	0.10	0.0082	40.4	0.20	0.0082	69.4	1.00	0.0000	138.3
u2152	0.90	0.0012	62.3	1.00	0.0000	101.0	1.00	0.0000	201.6
u2319	0.00	0.0283	93.7	0.00	0.0318	153.1	0.10	0.0100	323.5
pr2392	1.00	0.0000	80.6	1.00	0.0000	134.2	1.00	0.0000	267.8
Average	0.79	0.0041	34.8	0.87	0.0025	57.1	0.93	0.0008	113.2

TABLE V
RESULTS BY LKH-V10 AND #STD_α|S&R CASES ON VLSI DATA SET

Name	f^*	LKH-V10		$N_P=150$		$N_P=1000$			
		RPD	$t_r(s)$	RPD _B	$t_r(s)$	p_s	RPD	$t_r(s)$	T_{RUN}
pia3056	8258	0.0484	3682	0.0000	60.5	0.20	0.0170	470.2	193.1
dke3097	10539	0.0000	4404	0.0000	60.8	1.00	0.0000	394.5	141.5
lsn3119	9114	0.0000	4180	0.0000	64.8	1.00	0.0000	440.0	163.8
lta3140	9517	0.0000	4611	0.0000	57.8	1.00	0.0000	425.7	158.4
beg3293	9772	0.0000	6115	0.0000	62.4	1.00	0.0000	425.0	143.9
fdp3256	10008	0.0100	6559	0.0000	68.2	0.50	0.0050	489.5	174.3
dhb3386	11137	0.0000	5165	0.0000	66.9	1.00	0.0000	461.3	148.5
fjs3649	9272	0.0000	9112	0.0065	81.8	1.00	0.0000	593.5	196.3
fjr3672	9601	0.0000	7969	0.0077	77.6	1.00	0.0000	579.8	189.0
dlb3694	10959	0.0000	6275	0.0000	98.2	0.10	0.0082	673.0	212.7
ltb3729	11821	0.0000	7599	0.0000	87.4	1.00	0.0000	601.3	182.0
xqe3891	11995	0.0083	7174	0.0000	99.4	1.00	0.0000	653.0	200.1
xua3937	11239	0.0000	6870	0.0100	88.3	0.60	0.0036	590.0	185.2
dke3938	12503	0.0000	6654	0.0000	108.5	1.00	0.0000	654.0	190.3
dkc3954	12538	0.0000	6864	0.0000	99.5	1.00	0.0000	684.8	197.2
bgb4355	12723	0.0472	9647	0.0000	134.3	1.00	0.0000	855.5	221.2
bgd4396	13009	0.0000	9865	0.0000	111.9	1.00	0.0000	744.2	198.9
frv4410	10711	0.0000	9200	0.0000	115.3	0.90	0.0009	847.8	237.4
bgf4475	13221	0.0000	17943	0.0000	114.1	1.00	0.0000	765.6	196.3
xqd4966	15316	0.0000	15280	0.0000	177.4	1.00	0.0000	1118	255.9
fqm5087	13029	0.0000	30226	0.0000	198.4	0.70	0.0023	1229	267.2
fea5557	15445	0.0000	20002	0.0000	196.3	1.00	0.0000	1352	280.6
Average -		0.0051	9336	0.0011	101.4	0.86	0.0017	684.0	197.0

$N_P = 500$ achieves performance 46.92 times lower for the RPD and 16.64 times faster for the normalized running time than that of IBGLK; and the MAOS with $N_P = 1000$ achieves performance 1.01 times higher for the p_s , 2.63 times lower for the RPD, and 6.89 times faster for the normalized running time than that of PHGA.

Table V compares the results of LKH-V10 [7] and the #STD_α|S&R with $N_P = 150, 1000$ for the instances in VLSI data set, where RPD_B means the best RPD. LKH-V10 [7] is the best of ten runs of LKH, each with V iterations, on AMD Athlon 1900+ processor (1.6 GHz) [7]. The MAOS with $N_P = 1000$ achieves performance 3.00 times lower for the RPD and 13.65 times faster for the real running time than that

TABLE VI
RESULTS BY #STD_α|S&R AND TWO EAX VARIANTS
ON TSPLIB INSTANCES

Name	EAX-Dis			HeSEA			#STD _α S&R ($N_P=300$)		
	p_s	RPD	$t_r(s)$	p_s	RPD	$t_r(s)$	p_s	RPD	$t_r(s)$
pr1002	-	-	-	1.00	0.0000	91	1.00	0.0000	19.5
vm1084	0.94	0.0011	77	1.00	0.0000	80.6	1.00	0.0000	17.0
pcb1173	0.98	0.0000	70	1.00	0.0000	84.5	1.00	0.0000	21.4
u1432	0.52	0.0105	150	1.00	0.0000	107	0.80	0.0024	36.5
vm1748	1.00	0.0000	253	1.00	0.0000	141	1.00	0.0000	40.8
u2152	-	-	-	1.00	0.0000	211	0.90	0.0012	62.3
pr2392	1.00	0.0000	464	1.00	0.0000	208	1.00	0.0000	80.6

of LKH-V10. Moreover, LKH-V10 can be viewed as ten runs of LKH-V. Then, for LKH-V, its RPD_B is exactly the RPD of LKH-V10, and its running time is 1/10 of the t_r of LKH-V10. The MAOS with $N_P = 150$ achieves performance 4.64 times lower for the RPD_B and 9.21 times faster for the real running time than that of LKH-V.

Table VI compares the results of the #STD_α|S&R with $N_P = 300$ and two prominent EAX variants, i.e., EAX-Dis [37] and HeSEA [36]. Only HeSEA [36] integrates an advanced LS strategy, i.e., a CLK [16] with an ICG in family competition. The comparison is not complete since the two EAX variants were only tested on five and seven selected instances for the 21 TSPLIB instances, with $1000 \leq V < 3000$. For EAX-Dis, it is coded in C++ and was run on a 1.7-GHz Xeon processor. For HeSEA, it is coded in C and was executed on a 1.2-GHz Pentium IV processor. For the five selected TSP instances, #STD_α|S&R achieves performance 1.08 times higher for the p_s , 4.83 times lower for the RPD, and 5.17 times faster for the real running time than that of EAX-Dis [30]. For the seven selected TSP instances, #STD_α|S&R obtains slightly worse p_s than HeSEA on u1432 and u2152 and the same p_s ($= 1.00$) as HeSEA on other five instances, and achieves performance 4.67 times faster than HeSEA for the real running time.

Unlike many state-of-the-art metaheuristics [15], [27], [30], including some EAX variants, such as HeSEA [36], current MAOS versions do not use explicit LS strategies, particularly the LK variants, during the runtime. In fact, The MAOS cases are good at certain TSP instances, while LS strategies may be good at some other TSP instances. For example, LKH is good at u2319, u1432, pr1002, etc., while the MAOS cases are good at si1032, d1291, vm1748, etc.

There are two ways to implement LS strategies into MAOS. The first is to chain R_{XS} with R_{LS} , as in GLS. The second is to consider R_{LS} rules as alternate rules in the macro R_{XS} rule in switch mode, since R_{LS} can be seen as a special R_{XS} , which uses only one parent as the incumbent state.

Due to the machine memory limitation, the MAOS cases are evaluated on some TSP instances with $V < 6000$ only. For tackling large-scale instances, it may be interesting to incorporate MAOS with some edge-locking approaches [5], [81], which may reduce those instances into smaller ones by locking certain edges, e.g., the common segments in tours [30].

VI. CONCLUSION

The MAOS is organized with a group of agents self-organizing in the ENV, where they indirectly cooperate with

others through an IC. In the simplified MAOS version, each agent possesses extremely limited declarative knowledge and simple procedural rules. Moreover, using generic macro search rules to embed with some basic search rule is addressed. MAOS may be particularly suitable for supporting cooperative search.

The implementation of MAOS focuses on utilizing simple and efficient components for tailoring to the problem domain. Moreover, two simple improving heuristics are taken into account based on a GEAX.

The experimental results on both the TSPLIB instances and the VLSI data set demonstrate that MAOS is competitive with some state-of-the-art algorithms, including LKH, PHGA, and IBGLK, in both the solution quality and the running time.

The characteristics of the components and parameters of MAOS are investigated based on the overall performance. The advantage of two simple improving heuristics is also shown. In particular, with the macro rule in switch mode, which deploys alternate rules with offline performance in negative correlations, agents may achieve better performance than with each alternate rule alone. It may prevent the difficulty of finding an omnipotent rule for tackling a full set of tasks.

Future works may be performed on the following topics: 1) search for new alternate rules that are particularly good for some difficult problem instances; 2) study online strategies for adaptive deployment of alternate rules during the runtime; and 3) apply MAOS to other hard computational problems.

REFERENCES

- [1] D. S. Johnson and L. A. McGeoch, "Experimental analysis of heuristics for the STSP," in *The Traveling Salesman Problem and its Variations*, G. Gutin and A. Punnen, Eds. Norwell, MA: Kluwer, 2002, pp. 369–443. [Online]. Available: <http://www.research.att.com/~dsj/chtsp/>
- [2] G. Reinelt, *The Traveling Salesman: Computational Solutions for TSP Applications*. Berlin, Germany: Springer-Verlag, 1994.
- [3] J. Y. Lee, S. Y. Shin, T. H. Park, and B. T. Zhang, "Solving traveling salesman problems with DNA molecules encoding numerical values," *Biosystems*, vol. 78, no. 1–3, pp. 39–47, Dec. 2004.
- [4] R. Martoňák, G. E. Santoro, and E. Tosatti, "Quantum annealing of the traveling-salesman problem," *Phys. Rev. E, Stat. Phys. Plasmas Fluids Relat. Interdiscip. Top.*, vol. 70, no. 5, p. 057701, Nov. 2004.
- [5] C. Walshaw, "A multilevel approach to the travelling salesman problem," *Oper. Res.*, vol. 50, no. 5, pp. 862–877, Sep./Oct. 2002.
- [6] G. Zaránd, F. Pázmándi, K. F. Pál, and G. T. Zimányi, "Using hysteresis for optimization," *Phys. Rev. Lett.*, vol. 89, no. 15, p. 150201, Oct. 2002.
- [7] W. Cook, *VLSI Data Sets*, 2003. [Online]. Available: <http://www.tsp.gatech.edu/vlsi/>
- [8] S. Climer and W. Zhang, "Take a walk and cluster genes: A TSP-based approach to optimal rearrangement clustering," in *Proc. Int. Conf. Mach. Learn.*, Banff, AB, Canada, 2004, pp. 169–176.
- [9] O. Johnson and J. Liu, "A traveling salesman approach for predicting protein functions," *Source Code Biol. Med.*, vol. 1, pp. 1–7, 2006. Art. No. 3.
- [10] M. R. Garey and D. S. Johnson, *Computers and Intractability: A Guide to the Theory of NP-Completeness*. San Francisco, CA: Freeman, 1979.
- [11] G. Zweig, "An effective tour construction and improvement procedure for the traveling salesman problem," *Oper. Res.*, vol. 43, no. 6, pp. 1049–1057, Nov./Dec. 1995.
- [12] S. Lin and B. W. Kernighan, "An effective heuristic algorithm for the traveling-salesman problem," *Oper. Res.*, vol. 21, no. 2, pp. 498–516, Mar./Apr. 1973.
- [13] N. Christofides and S. Eilon, "Algorithms for large-scale travelling salesman problems," *Oper. Res. Q.*, vol. 23, no. 4, pp. 511–518, Dec. 1972.
- [14] C. P. Gomes and B. Selman, "Algorithm portfolios," *Artif. Intell.*, vol. 126, no. 1/2, pp. 43–62, Feb. 2001.
- [15] K. Helsgaun, "An effective implementation of the Lin–Kernighan traveling salesman heuristic," *Eur. J. Oper. Res.*, vol. 126, no. 1, pp. 106–130, 2000. [Online]. Available: <http://www.akira.ruc.dk/~keld/research/LKH/>
- [16] D. Applegate, W. Cook, and A. Rohe, "Chained Lin–Kernighan for large traveling salesman problems," *INFORMS J. Comput.*, vol. 15, no. 1, pp. 82–92, Jan. 2003.
- [17] O. C. Martin, S. W. Otto, and E. W. Felten, "Large-step Markov chains for the traveling salesman problem," *Complex Syst.*, vol. 5, no. 3, pp. 299–326, 1991.
- [18] T. Stützle, "Local search algorithms for combinatorial problems: Analysis, improvements, and new applications," Ph.D. dissertation, Dept. Comput. Sci., Darmstadt Univ. Technol., Darmstadt, Germany, 1999.
- [19] W. Zhang and M. Looks, "A novel local search algorithm for the traveling salesman problem that exploits backbones," in *Proc. Int. Joint Conf. Artif. Intell.*, Edinburgh, U.K., 2005, pp. 343–348.
- [20] J. H. Holland, *Adaptation in Natural and Artificial Systems*. Ann Arbor, MI: Univ. Michigan Press, 1975.
- [21] D. E. Goldberg and R. Lingle, "Alleles, loci, and the traveling salesman problem," in *Proc. Int. Conf. Genetic Algorithms Their Appl.*, Pittsburgh, PA, 1985, pp. 154–159.
- [22] H. Mühlenbein, "Evolution in time and space—The parallel genetic algorithm," in *Foundations of Genetic Algorithms*. San Mateo, CA: Morgan Kaufmann, 1991, pp. 316–337.
- [23] I. Oliver, D. Smith, and J. H. Holland, "A study of permutation crossover operators on the traveling salesman problem," in *Proc. Int. Conf. Genetic Algorithms*, 1987, pp. 224–230.
- [24] T. Guo and Z. Michalewicz, "Inver-over operator for the TSP," in *Proc. Int. Conf. Parallel Problem Solving From Nature*, Amsterdam, The Netherlands, 1998, pp. 803–812.
- [25] S. Jung and B. R. Moon, "The natural crossover for the 2D Euclidean TSP," in *Proc. Genetic Evol. Comput. Conf.*, Las Vegas, NV, 2000, pp. 1003–1010.
- [26] K. Maekawa, N. Mori, H. Tamaki, H. Kita, and Y. Nishikawa, "A genetic solution for the traveling salesman problem by means of a thermodynamical selection rule," in *Proc. IEEE Int. Conf. Evol. Comput.*, Nagoya, Japan, 1996, pp. 529–534.
- [27] P. Merz and B. Freisleben, "Memetic algorithms for the traveling salesman problem," *Complex Syst.*, vol. 13, no. 4, pp. 297–345, 2001.
- [28] Y. Nagata and S. Kobayashi, "Edge assembly crossover: A high-power genetic algorithm for the traveling salesman problem," in *Proc. Int. Conf. Genetic Algorithms*, East Lansing, MI, 1997, pp. 450–457.
- [29] D. Whitley, T. Starkweather, and D. A. Fuquay, "Scheduling problems and traveling salesman: The genetic edge recombination," in *Proc. Int. Conf. Genetic Algorithms*, Fairfax, VA, 1989, pp. 133–140.
- [30] H. D. Nguyen, I. Yoshihara, K. Yamamori, and M. Yasunaga, "Implementation of an effective hybrid GA for large-scale traveling salesman problems," *IEEE Trans. Syst., Man, Cybern. B, Cybern.*, vol. 37, no. 1, pp. 92–99, Feb. 2007.
- [31] M. Dorigo, V. Maniezzo, and A. Colnori, "Ant system: Optimization by a colony of cooperating agents," *IEEE Trans. Syst., Man, Cybern. B, Cybern.*, vol. 26, no. 1, pp. 29–41, Feb. 1996.
- [32] T. Stützle and H. H. Hoos, "MAX–MIN ant system," *Future Gener. Comput. Syst.*, vol. 16, no. 8, pp. 889–914, Jun. 2000.
- [33] J. Kennedy, R. C. Eberhart, and Y. Shi, *Swarm Intelligence*. San Mateo, CA: Morgan Kaufmann, 2001.
- [34] X. H. Shi, Y. C. Liang, H. P. Lee, C. Lu, and Q. X. Wang, "Particle swarm optimization-based algorithms for TSP and generalized TSP," *Inf. Process. Lett.*, vol. 103, no. 5, pp. 169–176, Aug. 2007.
- [35] R. Baraglia, J. I. Hidalgo, and R. Perego, "A hybrid heuristic for the traveling salesman problem," *IEEE Trans. Evol. Comput.*, vol. 5, no. 6, pp. 613–622, Dec. 2001.
- [36] H. K. Tsai, J. M. Yang, Y. F. Tsai, and C. Y. Kao, "An evolutionary algorithm for large traveling salesman problems," *IEEE Trans. Syst., Man, Cybern. B, Cybern.*, vol. 34, no. 4, pp. 1718–1729, Aug. 2004.
- [37] Y. Nagata, "The EAX algorithm considering diversity loss," in *Proc. Int. Conf. Parallel Problem Solving From Nature*, Birmingham, U.K., 2004, pp. 332–341.
- [38] X.-F. Xie and J. Liu, "A compact multiagent system based on autonomy oriented computing," in *Proc. IEEE/WIC/ACM Int. Conf. Intell. Agent Technol.*, Compiègne, France, 2005, pp. 38–44.
- [39] G. Di Marzo Serugendo, M.-P. Gleizes, and A. Karageorgos, "Self-organisation and emergence in MAS: An overview," *Informatica*, vol. 30, pp. 45–54, 2006.
- [40] A. M. Glenberg, "What memory is for," *Behav. Brain Sci.*, vol. 20, pp. 1–55, 1997.
- [41] G. Gigerenzer and D. G. Goldstein, "Reasoning the fast and frugal way: Models of bounded rationality," *Psychol. Rev.*, vol. 103, no. 4, pp. 650–669, Oct. 1996.
- [42] X.-F. Xie and W.-J. Zhang, "SWAF: Swarm algorithm framework for numerical optimization," in *Proc. Genetic Evol. Comput. Conf.*, Seattle, WA, 2004, pp. 238–250.

- [43] B. G. Galef, "Why behaviour patterns that animals learn socially are locally adaptive," *Anim. Behav.*, vol. 49, no. 5, pp. 1325–1334, May 1995.
- [44] L.-A. Giraldeau, T. Caraco, and T. J. Valone, "Social foraging: Individual learning and cultural transmission of innovations," *Behav. Ecol.*, vol. 5, no. 1, pp. 35–43, 1994.
- [45] É. Danchin, L.-A. Giraldeau, T. J. Valone, and R. H. Wagner, "Public information: From nosy neighbors to cultural evolution," *Science*, vol. 305, no. 5683, pp. 487–491, Jul. 2004.
- [46] J. Liu and K.-C. Tsui, "Toward nature-inspired computing," *Commun. ACM*, vol. 49, no. 10, pp. 59–64, Oct. 2006.
- [47] D. Weyns and T. Holvoet, "On the role of environments in multiagent systems," *Informatica*, vol. 29, pp. 409–421, 2005.
- [48] X.-F. Xie and J. Liu, "How autonomy oriented computing (AOC) tackles a computationally hard optimization problem," in *Proc. Int. Joint Conf. Auton. Agents Multiagent Syst.*, Hakodate, Japan, 2006, pp. 646–653.
- [49] J. R. Anderson, "Human symbol manipulation within an integrated cognitive architecture," *Cogn. Sci.*, vol. 29, no. 3, pp. 313–341, May 2005.
- [50] A. Newell and H. A. Simon, *Human Problem Solving*. Englewood Cliffs, NJ: Prentice-Hall, 1972.
- [51] L. Cardelli, "Type systems," *ACM Comput. Surv.*, vol. 28, no. 1, pp. 263–264, Mar. 1996.
- [52] E. D. Taillard, L. M. Gambardella, M. Gendreau, and J.-Y. Potvin, "Adaptive memory programming: A unified view of metaheuristics," *Eur. J. Oper. Res.*, vol. 135, no. 1, pp. 1–16, Nov. 2001.
- [53] T. Smith, P. Husbands, P. Layzell, and M. O'Shea, "Fitness landscapes and evolvability," *Evol. Comput.*, vol. 10, no. 1, pp. 1–34, 2002.
- [54] M. J. Wooldridge and N. R. Jennings, "Intelligent agents: Theory and practice," *Knowl. Eng. Rev.*, vol. 10, no. 2, pp. 115–152, 1995.
- [55] A. Bandura, *Social Foundations of Thought and Action: A Social Cognitive Theory*. Englewood Cliffs, NJ: Prentice-Hall, 1986.
- [56] R. Boyd and P. J. Richerson, *The Origin and Evolution of Cultures*. New York: Oxford Univ. Press, 2005.
- [57] M. V. Flinn, "Culture and the evolution of social learning," *Evol. Hum. Behav.*, vol. 18, no. 1, pp. 23–67, Jan. 1997.
- [58] D. Curran and C. O'Riordan, "Increasing population diversity through cultural learning," *Adapt. Behav.*, vol. 14, no. 4, pp. 315–338, Dec. 2006.
- [59] A. Baddeley, "Exploring the central executive," *Q. J. Exp. Psychol., A*, vol. 49, no. 1, pp. 5–28, Feb. 1996.
- [60] T. G. Dietterich, "Learning at the knowledge level," *Mach. Learn.*, vol. 1, no. 3, pp. 287–315, Sep. 1986.
- [61] A. K. Romney, J. P. Boyd, C. C. Moore, W. H. Batchelder, and T. J. Brazill, "Culture as shared cognitive representations," *Proc. Nat. Acad. Sci. U.S.A.*, vol. 93, no. 10, pp. 4699–4705, May 1996.
- [62] I. D. Couzin, J. Krause, R. James, G. D. Ruxton, and N. R. Franks, "Collective memory and spatial sorting in animal groups," *J. Theor. Biol.*, vol. 218, no. 1, pp. 1–11, Sep. 2002.
- [63] J. K. Olick and J. Robbins, "Social memory studies: From "collective memory" to the historical sociology of mnemonic practices," *Annu. Rev. Sociology*, vol. 24, pp. 105–140, 1998.
- [64] T. Edgington, B. Choi, K. Henson, T. S. Raghu, and A. Vinze, "Adopting ontology to facilitate knowledge sharing," *Commun. ACM*, vol. 47, no. 11, pp. 85–90, Nov. 2004.
- [65] P. Stone and M. Veloso, "Multiagent systems: A survey from a machine learning perspective," *Auton. Robots*, vol. 8, no. 3, pp. 345–383, Jun. 2000.
- [66] E. Bonabeau, "Agent-based modeling: Methods and techniques for simulating human systems," *Proc. Nat. Acad. Sci. U.S.A.*, vol. 99, no. 10, pp. 7280–7287, May 2002.
- [67] B. A. Huberman, R. M. Lukose, and T. Hogg, "An economics approach to hard computational problems," *Science*, vol. 275, no. 5296, pp. 51–54, Jan. 1997.
- [68] M. Milano and A. Roli, "MAGMA: A multiagent architecture for metaheuristics," *IEEE Trans. Syst., Man, Cybern. B, Cybern.*, vol. 34, no. 2, pp. 925–941, Apr. 2004.
- [69] M. Toulouse, T. G. Crainic, and B. Sanso, "Systemic behavior of cooperative search algorithms," *Parallel Comput.*, vol. 30, no. 1, pp. 57–79, Jan. 2004.
- [70] K. D. Boese, A. B. Kahng, and S. Muddu, "A new adaptive multi-start technique for combinatorial global optimizations," *Oper. Res. Lett.*, vol. 16, no. 2, pp. 101–113, Sep. 1994.
- [71] D. Gamboa, C. Rego, and F. Glover, "Data structures and ejection chains for solving large-scale traveling salesman problems," *Eur. J. Oper. Res.*, vol. 160, no. 1, pp. 154–171, Jan. 2005.
- [72] M. Held and R. M. Karp, "The traveling-salesman problem and minimum spanning trees," *Oper. Res.*, vol. 18, no. 6, pp. 1138–1162, Nov./Dec. 1970.
- [73] J. L. Bentley, "Fast algorithms for geometric traveling salesman problems," *ORSA J. Comput.*, vol. 4, no. 4, pp. 387–411, 1992.
- [74] W. Cook and P. Seymour, "Tour merging via branch-decomposition," *INFORMS J. Comput.*, vol. 15, no. 3, pp. 233–248, Jul. 2003.
- [75] J.-P. Watson, C. Ross, V. Eisele, J. Denton, J. Bins, C. Guerra, L. D. Whitley, and A. E. Howe, "The traveling salesrep problem, edge assembly crossover, and 2-opt," in *Proc. Int. Conf. Parallel Problem Solving From Nature*, Amsterdam, The Netherlands, 1998, pp. 823–832.
- [76] K. Ikeda and S. Kobayashi, "Deterministic multi-step crossover fusion: A handy crossover composition for GAs," in *Proc. Int. Conf. Parallel Problem Solving From Nature*, Granada, Spain, 2002, pp. 162–171.
- [77] M. J. Streeter, D. Golovin, and S. F. Smith, "Combining multiple heuristics online," in *Proc. Nat. Conf. Artif. Intell.*, Vancouver, BC, Canada, 2007, pp. 1197–1203.
- [78] P. Kilby, J. K. Slaney, and T. Walsh, "The backbone of the travelling salesperson," in *Proc. Int. Joint Conf. Artif. Intell.*, Edinburgh, U.K., 2005, pp. 175–180.
- [79] M. L. Fredman, D. S. Johnson, L. A. McGeoch, and G. Ostheimer, "Data structures for traveling salesmen," *J. Algorithms*, vol. 18, pp. 432–479, 1995.
- [80] G. Reinelt, "TSPLIB—A traveling salesman problem library," *ORSA J. Comput.*, vol. 3, pp. 376–384, 1991.
- [81] T. Fischer and P. Merz, "Reducing the size of traveling salesman problem instances by fixing edges," in *Proc. Eur. Conf. Evol. Comput. Combinatorial Optimization*, Valencia, Spain, 2007, pp. 72–83.